

Practical: basic stats for peak calling

Jacques van Helden

2015-10-01

Peak-calling: Carl's question

What we want to do



Defining the data directory

Open a connection to the lab web browser.

We will first define the URL from which the data can be downloaded, by concatenating the URL fo the course with the path to our dataset.

To concatenate paths, it is *recommended* to use the **R** command `file.path()`.

```
url.course <- "http://jvanheld.github.io/EBA15_stats_RStudio"
url.data <- file.path(url.course, "practicals", "02_peak-ca
```

Loading a data table

R enables to download data directly from the Web.

Load counts per window in chip sample.

```
## Define URL of the ChIP file
chip.bedg.file <- file.path(url.data, "FNR_200bp.bedg")

## Load the file content in an R data.frame
chip.bedg <- read.table(chip.bedg.file)

## Set column names
names(chip.bedg) <- c("chrom", "start", "end", "counts")
```

Exploring a data frame: dim()

Before anything else, let us inspect the size of the data frame, in order to check that it was properly loaded.

```
dim(chip.bedg)
```

```
## [1] 23199      4
```

Checking the n first rows: head()

The function head() displays the first rows of a table.

```
head(chip.bedg, n = 5)
```

##		chrom	start	end	counts
## 1	gi 49175990 ref NC_000913.2		0	200	1594
## 2	gi 49175990 ref NC_000913.2		200	400	834
## 3	gi 49175990 ref NC_000913.2		400	600	222
## 4	gi 49175990 ref NC_000913.2		600	800	172
## 5	gi 49175990 ref NC_000913.2		800	1000	123

Checking the n last rows: tail()

The function `tail()` displays the first rows of a table.

```
head(chip.bedg, n = 5)
```

##		chrom	start	end	counts
## 1	gi 49175990 ref NC_000913.2		0	200	1594
## 2	gi 49175990 ref NC_000913.2		200	400	834
## 3	gi 49175990 ref NC_000913.2		400	600	222
## 4	gi 49175990 ref NC_000913.2		600	800	172
## 5	gi 49175990 ref NC_000913.2		800	1000	123

Viewing a table

The function `View()` displays the full table in a user-friendly mode.

```
View(chip.bedg)
```


Selecting arbitrary rows

```
chip.bedg[100:105,]
```

##		chrom	start	end	counts
## 100	gi 49175990 ref NC_000913.2	19800	20000	21	
## 101	gi 49175990 ref NC_000913.2	20000	20200	0	
## 102	gi 49175990 ref NC_000913.2	20200	20400	0	
## 103	gi 49175990 ref NC_000913.2	20400	20600	108	
## 104	gi 49175990 ref NC_000913.2	20600	20800	229	
## 105	gi 49175990 ref NC_000913.2	20800	21000	245	

Selecting arbitrary columns

```
chip.bedg[100:105, 2]
```

```
## [1] 19800 20000 20200 20400 20600 20800
```

```
chip.bedg[100:105, "start"]
```

```
## [1] 19800 20000 20200 20400 20600 20800
```

```
chip.bedg[100:105, c("start", "counts")]
```

```
##      start counts
## 100 19800      21
## 101 20000       0
## 102 20200       0
## 103 20400     108
## 104 20600     229
## 105 20800     245
```

Adding columns

We can add columns with the result of computations from other columns.


```
chip.bedg$midpos <- (chip.bedg$start + chip.bedg$end)/2  
head(chip.bedg)
```

##		chrom	start	end	counts	midpos
## 1	gi 49175990 ref NC_000913.2		0	200	1594	100
## 2	gi 49175990 ref NC_000913.2		200	400	834	300
## 3	gi 49175990 ref NC_000913.2		400	600	222	500
## 4	gi 49175990 ref NC_000913.2		600	800	172	700
## 5	gi 49175990 ref NC_000913.2		800	1000	123	900
## 6	gi 49175990 ref NC_000913.2		1000	1200	116	1100

Plotting a density profile

We can readily print a plot with the counts per window.


```
plot(chip.bedg[, c("midpos", "counts")], type="h")
```



Plotting a density profile

Let us improve the plot

```
plot(chip.bedg[, c("midpos", "counts")], type="h",  
     col="green", xlab="Genomic position (200bp windows)",  
     ylab= "Read counts",  
     main="FNR ChIP-seq")
```



Exercise: exploring the background

We already loaded the count table for the FNR ChIP counts per window.

The background level will be estimated by loading counts per window in a genomic input sample. These counts are available in the same directory a file named `input_200bp.bedg`

1. Load the counts per window in the input sample (genome sequencing).
2. Plot the density profile of the input
3. Compare chip-seq and input density profiles
4. Compare counts per window between chip-seq and input

Solution: loading the input counts per window

```
## Define URL of the input file
input.bedg.file <- file.path(url.data, "input_200bp.bedg")

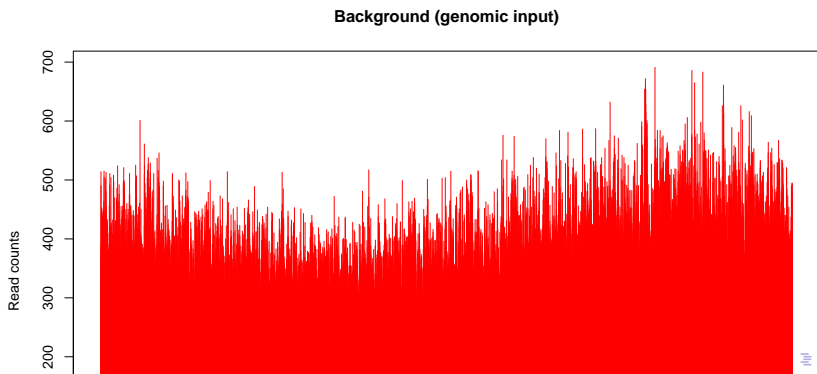
## Load the file content in an R data.frame
input.bedg <- read.table(input.bedg.file)

## Set column names
names(input.bedg) <- c("chrom", "start", "end", "counts")
```

Solution: plotting the input density profile


```
## Compute middle positions per window
input.bedg$midpos <- (input.bedg$start + input.bedg$end)/2

plot(input.bedg[, c("midpos", "counts")], type="h",
     col="red", xlab="Genomic position (200bp windows)",
     ylab= "Read counts",
     main="Background (genomic input)")
```



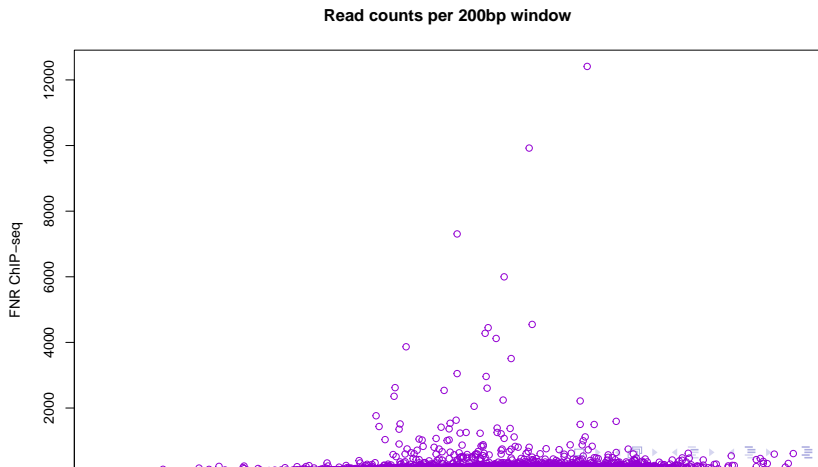
Solution: comparing chip-seq and background density profiles

```
par(mfrow=c(2,1)) ## Draw two panels on top of each other
plot(chip.bedg[, c("midpos", "counts")], type="h",
     col="green", xlab="Genomic position (200bp windows)",
     ylab= "Read counts",
     main="FNR ChIP-seq")
plot(input.bedg[, c("midpos", "counts")], type="h",
     col="red", xlab="Genomic position (200bp windows)",
     ylab= "Read counts",
     main="Background (genomic input)")
```



Solution: comparing counts per window between chip-seq and input

```
plot(input.bedg$counts, chip.bedg$counts, col="darkviolet",  
      xlab="Genomic input", ylab="FNR ChIP-seq",  
      main="Read counts per 200bp window")
```



Solution: comparing counts per window between chip-seq and input

In order to better highlight the dynamic range, we can use a log-based representation

```
plot(input.bedg$counts, chip.bedg$counts, col="darkviolet",  
      xlab="Genomic input", ylab="FNR ChIP-seq",  
      main="Read counts per 200bp window",  
      log="xy")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 377 x v  
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 403 y v  
## from logarithmic plot
```

```
grid() ## add a grid
```

Read counts per 200bp window



Questions

- ▶ On the ChIP-seq versus input plot, how would you define peaks ?
- ▶ Where would you place the limit between peaks and background fluctuations ?

Exercises

1. Think about further drawing modes to improve your perception of the differences between signal and background.
2. We will formulate (together) a reasoning path to compute a p-value for each peak.