

Exploration d'une table d'annotations génomiques (GTF)

Ecole de Bioinformatique Aviesan-IFB 2018

Jacques van Helden

2018-11-27

Contents

But de ce TP	1
Le format GTF	2
Localiser l'URL d'un fichier GTF	2
Page d'accueil d'EnsemblGenomes	2
EnsemblGenomes Fungi	3
EnsemblGenomes Download page	3
Le chemin de la maison (automatique)	4
Créer un espace de travail	4
Downloading the GTF file	5
Downloading the GTF file: solution	5
Loading a data table in R	5
Loading the GTF file	6
Exploring the content of a data table	6
Checking heads and tails	6
Viewing a table	6
Selecting columns	7
Feature types	7
Décompte par valeur	7
Computing feature lengths	9
Filtering rows based on a column content	9
Selecting genes from the GTF table	9
Downloading chromosome sizes	9
Loading chromosome sizes	10
Exercices	10
Gene length histogram	10
Setting a relevant number of breaks	10
Gene length distribution – improving the output	11
Getting the hist() data	11
Gene length box plot	12

But de ce TP

Durant ce TP, vous serez amenés à effectuer les tâches suivantes:

1. Manipuler une table de données génomique (les annotations du génome de la levure).
2. Sélectionner un sous-ensemble des données en filtrant les lignes sur base d'un critère déterminé (type d'annotation, chromosome).
3. Générer des graphiques pour représenter différents aspects liés à ces données.
4. Calculer des statistiques qui résument les différents types d'annotations.

Le format GTF

Le format **GTF (General Transfer Format)** est très largement utilisé pour fournir des annotations génomiques dans un format facilement lisible, tout en étant facilement manipulable au moyen de l'ordinateur.

Fichiers textuels,

- une ligne par "objet" génomique (gène, transcrit, exon, intron, CDS, ...)
- une colonne par attribut (nom, source, type d'objet, coordonnées génomique, description).

Le format est décrit sur les sites suivants.

- <http://www.ensembl.org/info/website/upload/gff.html>
- <https://genome.ucsc.edu/FAQ/FAQformat.html#format4>

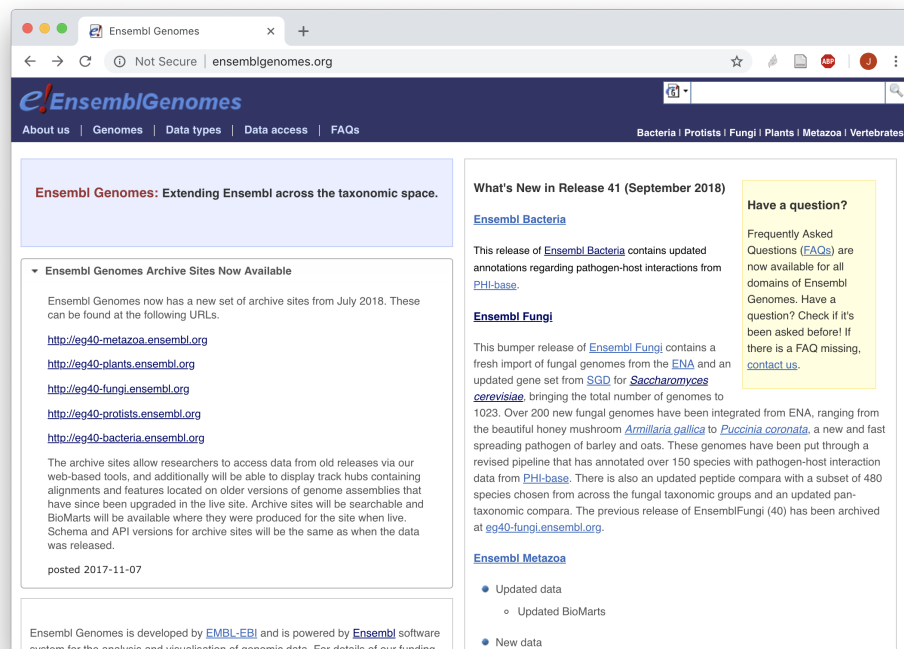
Localiser l'URL d'un fichier GTF

N'hésitez pas à adapter le protocole ci-dessous pour travailler avec votre propre génome.

1. Connectez-vous à <http://ensemblgenomes.org/>.
2. Cliquez sur le lien Fungi.
3. Cliquez Download
4. Dans la boîte **Filter**, tapez *Saccharomyces cerevisiae*. Pendant que vous écrivez, la liste des organismes proposés s'affine.
5. Copiez le lien du fichier gtf (*Saccharomyces_cerevisiae*.R64-1-1.41.gtf.gz).

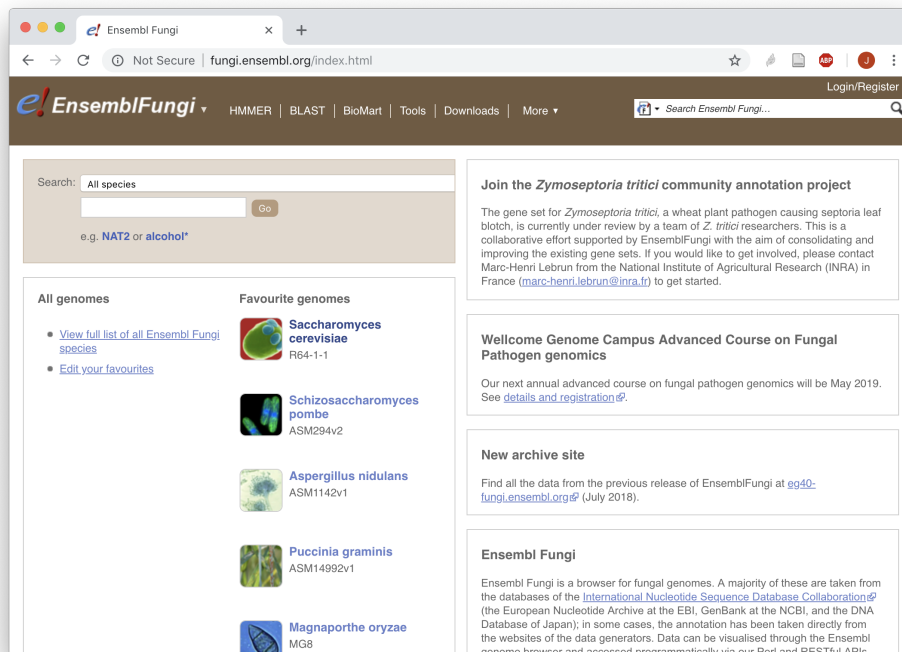
Page d'accueil d'EnsemblGenomes

```
include_graphics(path = "images/ensemblgenomes_home.png")
```



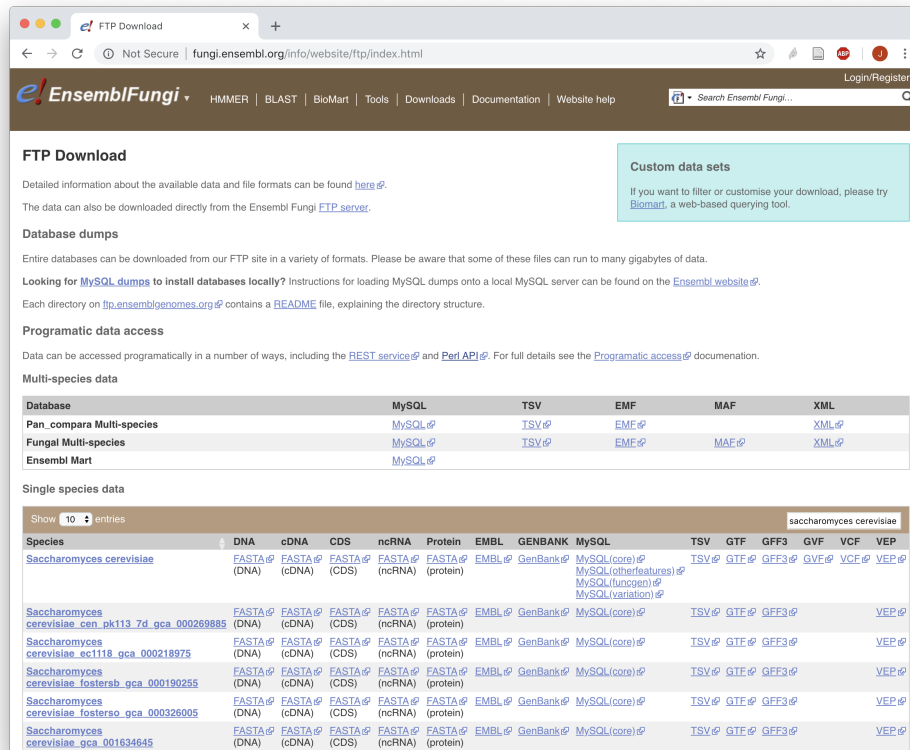
EnsemblGenomes Fungi

```
include_graphics(path = "images/ensemblgenomes_fungi.png")
```



EnsemblGenomes Download page

```
include_graphics(path = "images/ensemblgenomes_download_yeast.png")
```



Le chemin de la maison (automatique)

Sous Linux et Mac OS X, on peut identifier la racine de son compte avec la commande `R Sys.getenv()`.

- Invoquée sans paramètre, cette commande liste toutes les variables d'environnement (votre configuration système).
- On peut restreindre l'output à une variable d'environnement donnée, par exemple `Sys.getenv("HOME")` retourne le chemin de la racine de votre compte.
- Une écriture équivalente : le symbole tilde `~` indique également le chemin de la racine de votre compte.
- La notation ``~`` fonctionne également sous Windows, nous l'utiliserons donc ci-dessous.

Créer un espace de travail

Exercice: créer un dossier de travail nommé `workDir` à la racine de votre compte, et déplacez-vous dans ce dossier.

Solution ci-dessous.

```
## Define the working directory
workDir <- "~/intro_R/explorer_un_GTF"

## Create the working directory
dir.create(workDir, recursive = TRUE, showWarnings = FALSE)

## Go to the working directory
setwd(workDir)
getwd()      ## Check your current location
```

```
[1] "/Users/jvanheld/intro_R/explorer_un_GTF"
list.files() ## List files (should be empty if just created)
```

```
[1] "chrom_sizes.tsv"
[2] "Saccharomyces_cerevisiae.R64-1-1.41.gtf.gz"
```

Downloading the GTF file

Exercise: download the GTF file in the working directory (optionally, adapt the command to load a GTF of your interest). Before downloading the file we check if it is already present in the working directory. If yes, we skip the download.

****Tip:**** use the commands `file.exists`, `download.file`.

Downloading the GTF file: solution

```
## Define the file name (without path) in a separate variable, we will need it later
gtf.file <- 'Saccharomyces_cerevisiae.R64-1-1.41.gtf.gz'

## Define the URL by concatenating the URL of the directory and the file name
gtf.url <- file.path('ftp://ftp.ensemblgenomes.org/pub/release-41/fungi/gtf/saccharomyces_cerevisiae/',

## Download the file, but only if not yet there
if (file.exists(gtf.file)) {
  message("GTF annotation file already there: ", gtf.file)
} else {
  message("Downloading GTF annotation file")
  download.file(url = gtf.url, destfile = gtf.file)
}

## Check the files in the work directory
list.files()
```

```
[1] "data"
[2] "figures"
[3] "gtf_exploration_files"
[4] "gtf_exploration.html"
[5] "gtf_exploration.md"
[6] "gtf_exploration.pdf"
[7] "gtf_exploration.Rmd"
[8] "images"
[9] "Saccharomyces_cerevisiae.R64-1-1.41.gtf.gz"
```

Loading a data table in R

Commands: `read.table`, `read.delim`, `read.csv`.

R includes several types of tabular structures (matrix, data.frame, table). The most widely used is `data.frame()`, which consists in a table of values with a type (strings, integer, ..) attached to each column, and names associated to rows and columns.

The function `read.table()` enables to read a text file containing tabular data, and to store its content in a variable.

Several functions derived from `read.table()` facilitate the loading of different formats.

- `read.delim()` for files where a particular character is used as column separator (by default the tab character “`\t`”).
- `read.csv()` for “comma-separated values” values.

Loading the GTF file

Load the GTF file in a variable named `featureTable`.

****Tip:**** command `read.delim`.

```
## Load GTF file in a data.frame
featureTable <- read.delim(
  gtf.file, comment.char = "#", sep="\t",
  header=FALSE, row.names = NULL)

## The GTF format has no header, but we can define it based on the specification
names(featureTable) <- c("seqname", "source", "feature", "start", "end", "score", "strand", "frame", "a
```

Exploring the content of a data table

Immediately after having loaded a data table, check its dimensions.

```
dim(featureTable) ## Dimensions of the table
```

```
[1] 41606      9
```

```
nrow(featureTable) ## Number of rows
```

```
[1] 41606
```

```
ncol(featureTable) ## Number of columns
```

```
[1] 9
```

Checking heads and tails

Displaying the full annotation table would not be very convenient, since it contains tens of thousands of rows.

We can display the first rows of the file with the function `head()`, and the last rows with `tail()`.

```
## Display the 5 first rows of the feature table
head(featureTable, n = 5)
```

```
## Display the 5 last rows of the feature table
tail(featureTable, n = 5)
```

Viewing a table

If you are using the **RStudio** environment, you can display the table in a dynamic viewer pane with the function `View()`.

```
## In RStudio, display the table in a separate tab
View(featureTable)
```

The `View()` function is interactive, so it should not be used in a script because it would perturbate its execution.

Selecting columns

The last column of GTF files is particularly heavy, it contains a lot of semi-structured information.

We can select the 8 first columns and display the 5 first rows of this sub-table.

```
## Column selection + head
head(featureTable[,1:8], n=5)
```

	seqname	source	feature	start	end	score	strand	frame
1	IV	sgd	gene	1802	2953	.	+	.
2	IV	sgd	transcript	1802	2953	.	+	.
3	IV	sgd	exon	1802	2953	.	+	.
4	IV	sgd	CDS	1802	2950	.	+	0
5	IV	sgd	start_codon	1802	1804	.	+	0

```
## Equivalent: selecting subsets of rows and columns
featureTable[1:5, 1:8]
```

	seqname	source	feature	start	end	score	strand	frame
1	IV	sgd	gene	1802	2953	.	+	.
2	IV	sgd	transcript	1802	2953	.	+	.
3	IV	sgd	exon	1802	2953	.	+	.
4	IV	sgd	CDS	1802	2950	.	+	0
5	IV	sgd	start_codon	1802	1804	.	+	0

Feature types

Exercise: the column *feature* of the GTF indicates the feature table.

- List the feature types found in the GTF
- Count the number of features per type, and sort them by decreasing values.

Tip: commands `unique`, `table` and `sort`.

```
## List the types of features
unique(featureTable$feature)
```

```
[1] gene          transcript     exon          CDS
[5] start_codon   stop_codon    five_prime_utr
7 Levels: CDS exon five_prime_utr gene start_codon ... transcript
```

```
## Count the number of features per type
sort(table(featureTable$feature), decreasing = TRUE)
```

exon	gene	transcript	CDS	start_codon
7416	7036	7036	6913	6601
stop_codon	five_prime_utr			
6600	4			

Décompte par valeur

La fonction `table()` permet de compter le nombre d'occurrences de chaque valeur dans un vecteur ou un tableau. Quelques exemples d'utilisation ci-dessous.

```
## Count the number of features per chromosome
table(featureTable$seqname)
```

I	II	III	IV	IX	Mito	V	VI	VII	VIII	X	XI	XII	XIII	XIV
731	2841	1170	5185	1520	312	2055	898	3688	2012	2511	2180	3690	3196	2712
XV	XVI													
3706	3199													

```
## Count the number of features per type
table(featureTable$feature)
```

CDS	exon	five_prime_utr	gene	start_codon
6913	7416	4	7036	6601
stop_codon	transcript			
6600	7036			

On peut calculer des tables de contingence en comptant le nombre de combinaisons entre 2 vecteurs (ou 2 colonnes d'un tableau).

```
## Table with two vectors
table(featureTable$feature, featureTable$seqname)
```

	I	II	III	IV	IX	Mito	V	VI	VII	VIII	X	XI	XII
CDS	120	483	192	870	252	59	338	146	605	340	412	361	604
exon	129	500	210	907	269	87	367	166	659	358	446	385	668
five_prime_utr	0	0	0	0	1	0	0	0	0	0	1	0	0
gene	124	472	200	868	258	55	352	154	629	336	428	369	631
start_codon	117	458	184	836	241	28	323	139	583	321	398	348	578
stop_codon	117	456	184	836	241	28	323	139	583	321	398	348	578
transcript	124	472	200	868	258	55	352	154	629	336	428	369	631

	XIII	XIV	XV	XVI
CDS	531	454	609	537
exon	573	478	647	567
five_prime_utr	1	0	1	0
gene	541	455	628	536
start_codon	504	435	596	512
stop_codon	505	435	597	511
transcript	541	455	628	536

```
## Same result with a 2-column data frame
table(featureTable[, c("feature", "seqname")])
```

	seqname												
feature	I	II	III	IV	IX	Mito	V	VI	VII	VIII	X	XI	XII
CDS	120	483	192	870	252	59	338	146	605	340	412	361	604
exon	129	500	210	907	269	87	367	166	659	358	446	385	668
five_prime_utr	0	0	0	0	1	0	0	0	0	0	1	0	0
gene	124	472	200	868	258	55	352	154	629	336	428	369	631
start_codon	117	458	184	836	241	28	323	139	583	321	398	348	578
stop_codon	117	456	184	836	241	28	323	139	583	321	398	348	578
transcript	124	472	200	868	258	55	352	154	629	336	428	369	631

	seqname				
feature	XIII	XIV	XV	XVI	
CDS	531	454	609	537	
exon	573	478	647	567	
five_prime_utr	1	0	1	0	
gene	541	455	628	536	


```

start_codon    504 435 596 512
stop_codon     505 435 597 511
transcript     541 455 628 536

```

Computing feature lengths

- Add a column with feature lengths.

Note about feature length computation (explain why) :

$$L = \text{end} - \text{start} + 1$$

```

## Add a column to the table with genes lengths
featureTable$length <- featureTable$end - featureTable$start + 1

```

Filtering rows based on a column content

The function `subset()` enables to select a subset of rows based on a filter applied to the content of one or several columns.

We can use it to select the subset of features corresponding to genes.

Selecting genes from the GTF table

- Select of genes from the GTF table and store them in a separate variable named `genes`.
- Compute summary statistics about gene lengths

Tip: commands `subset`, `summary`.

```

## Select subset of features having "CDS" as "feature" attribute
genes <- subset(featureTable, feature == "gene")

## Print a message with the number of genes
message("Number of genes: ", nrow(genes))

## Compute basic statistics on genes lengths
summary(genes$length)

```

```

Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  51     468     1005    1275    1717    14733

```

Downloading chromosome sizes

- Download chromosome sizes (`chrom_sizes.tsv`)

```

## Download tab-delimited file with chromosome sizes (unless already there)
chromsizes.url <- "https://github.com/jvanheld/stats_avec_RStudio_EBA/blob/gh-pages/practicals/gtf_expl
chrom.size.file <- file.path(workDir, "chrom_sizes.tsv")

if (file.exists(chrom.size.file)) {
} else {
  download.file(chromsizes.url, destfile = chrom.size.file)
}

```

NULL

Loading chromosome sizes

```
## Read chromosome sizes
chrom.size <- read.delim(
  file = chrom.size.file,
  header = FALSE, row.names = 1)

## Assign a name to the columns
names(chrom.size) <- c("chromID", "size")
# View(chrom.size)

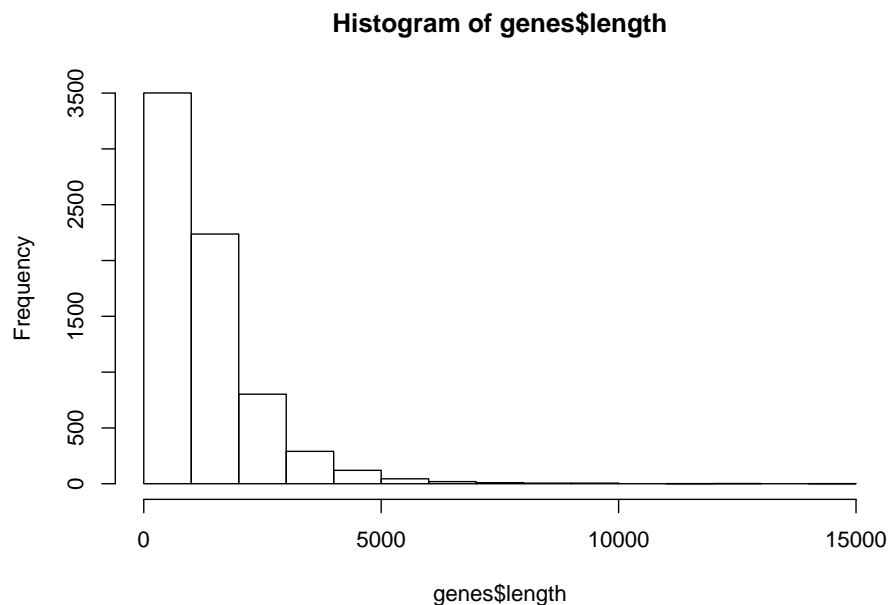
## print the size of the third chromosome
message("Length of chromosome III = ", chrom.size["III", "size"], " bp.")
```

Exercises

1. Draw a histogram with gene length distribution. Choose a relevant number of breaks to display an informative histogram.
2. Draw a barplot showing gene density per chromosome (number of genes per Mb).
3. Draw a boxplot of gene lengths per chromosome.

Gene length histogram

```
hist(genes$length)
```



Setting a relevant number of breaks

```
## Take more or less 100 bins
h <- hist(genes$length, breaks = 100)
```





Figure 1: Distribution of cds lengths for *Saccharomyces cerevisiae*.



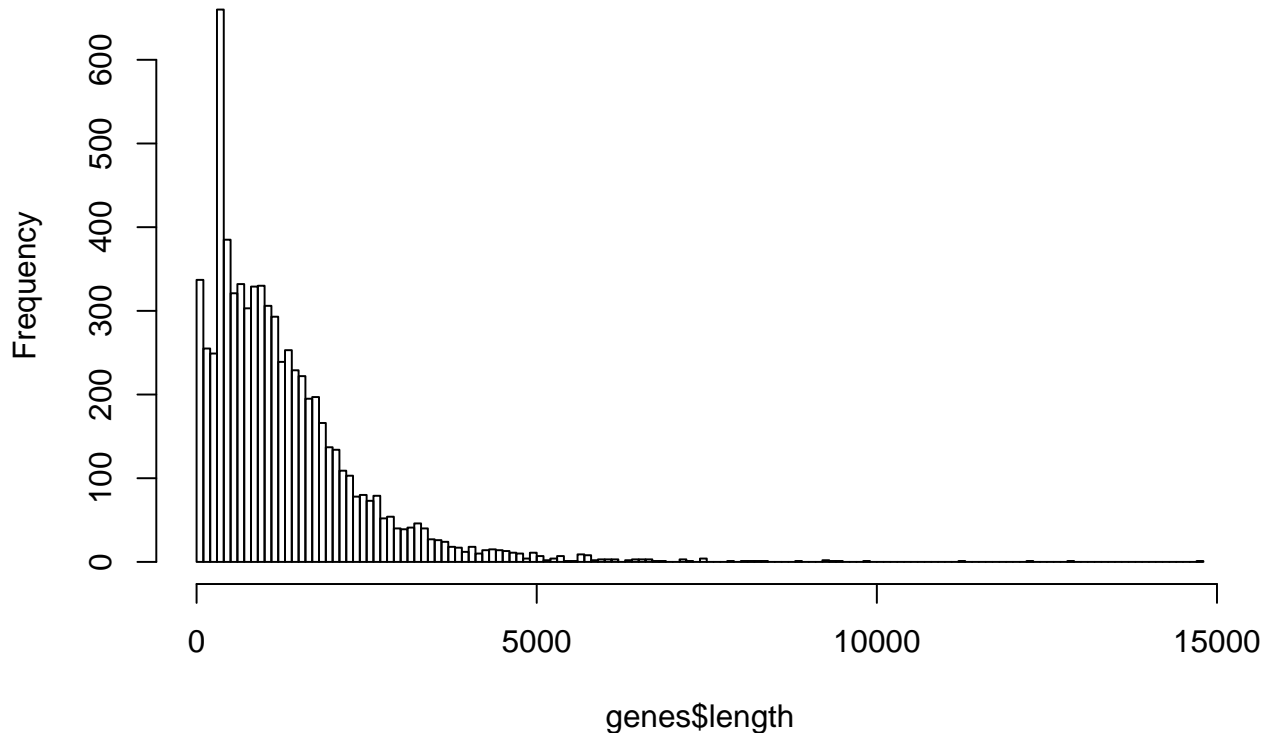
Gene length distribution – improving the output

Getting the hist() data

Récupérez le résultat de `hist()` dans une variable nommée `{histData}`.

```
## Define breaks exactly in the way you wish
histData <- hist(genes$length, breaks=seq(from=0, to=max(genes$length)+100, by=100))
```

Histogram of genes\$length



Imprimez le résultat à l'écran (`print()`) et analysez la structure de la variable `histData` (il s'agit d'une variable de type liste).

Fonctions utiles:

- `class(histData)`
- `attributes(histData)`
- `print(histData)`

```
## Display the values used to draw the histogram  
print(histData )
```

Gene length box plot

D'autres types de graphiques permettent d'explorer la distribution d'un ensemble des données. En particulier, les boîtes à moustaches (box plots) affichent, pour une série de données, la médiane, l'écart interquartile, un intervalle de confiance et les valeurs aberrantes.

```
boxplot(length ~ seqname, data = genes, col="palegreen", horizontal=TRUE, las=1, xlab="Gene length", ylab="length")
```




Figure 2: Boxplot of gene lengths per chromosome